

Anthony Tang

512-348-9289 | anthony.tang.tx@gmail.com | [LinkedIn](#) | [Website](#) | [Github](#) | Austin, TX

Dear Hiring Team,

I'm a senior at Rice University and I'm pursuing a bachelor's degree in Computer Science at Rice University with a 3.81 GPA, graduating in May 2027. I was born in Round Rock, Texas, but grew up in Austin, Texas, for most of my life. My coursework in Computer Systems, Concurrent Program Design, and Data Structures and Algorithms has also given me a strong foundation in the systems-level thinking that production engineering requires. I've built across the full stack with Python, TypeScript, React, Node.js, and PostgreSQL, and I work daily with Docker, Kubernetes, AWS, and CI/CD pipelines.

The best way to describe how I work is through what I did this summer at Nextdoor. I was on the Core Services team, where I designed and shipped the telemetry and monitoring system for their Streaming ORM. That meant instrumenting 8+ metric dimensions across production hot-path code and redesigning the team's Datadog dashboards from scratch so on-call engineers could actually diagnose problems in real time. The work required making careful decisions about what to measure and how to instrument without degrading latency on the hot path. It was the kind of project where getting it wrong meant either missing real issues or drowning the team in noise, so precision mattered.

Partway through that project, our observability costs started spiking. I dug into the data, traced the root cause to a regex bug producing cardinality explosions in our metric tags, fixed the bug, removed unnecessary tags, and staged the rollout carefully, monitoring each change in production before expanding it. The result was a 50% reduction in observability costs. I also designed an error-capture system for the same latency-sensitive hot-path code, owning the software design from architecture doc through implementation. By the end of the internship, I had won Best in Show at the company hackathon for building a new-advertiser onboarding tool with my team.

Before Nextdoor, I built production CI/CD infrastructure at yAI with automated testing in Python, achieving 99.9% deployment reliability. I also led a large-scale refactor that migrated 70% of the codebase to a modular microservices architecture. At Curriculum Associates, where I interned twice, I designed a high-performance query engine that processed 100,000+ data points and built Python automation tooling that eliminated 5+ hours of manual overhead per sprint. I won first-place hackathons at both companies. Across all four internships, the pattern has been the same: I ship production code, I debug hard problems, and I take ownership of the systems I build from design through deployment.

On the research side, I published a paper at the AAAI-26 conference where I trained logistic regression probes on the internal activations of Gemma-2 2B using PyTorch and Scikit-learn to test whether linear interpretability methods can predict when a large language model will fail. The probes returned near-chance performance across every layer and task category. The result contradicted my starting hypothesis, but the value of that project was in the rigor of the experimental design and in following the data to an honest conclusion rather than forcing the narrative I expected. That research instinct, building something carefully, testing it honestly, and iterating based on what the results actually say, carries over directly into how I write and debug production software.

Best,

Anthony Tang